

Scalable Generative Reconstruction of Three-Dimensional Fields on Point Clouds

Anonymous authors
Paper under double-blind review

Abstract

Reconstructing three-dimensional physical fields from sparse measurements is an ill-posed inverse problem for which generative models provide the natural formulation: they sample from the distribution over fields consistent with the data rather than regressing to its mean. Current methods do not survive the transition to three dimensions—voxel diffusion computes globally on dense grids and hits a measured memory wall, while latent methods compress fields through an autoencoder that bounds what can be recovered. We introduce LOCUS, a conditional rectified-flow model that generates fields directly in ambient function space, on arbitrary point clouds, with no autoencoder and no voxelization. Each query point is conditioned through a compact global summary of the observations and a local, importance-weighted retrieval over sensor tokens; the training objective, an integral over the domain, is estimated on a random $\sim 2\%$ subset of points per step, decoupling training cost from resolution. A single model is amortized over sensor counts, layouts, noise levels, and observed-channel patterns. On spatially disjoint isotropic-turbulence evaluation at two million points, LOCUS matches or beats a strong latent flow-matching baseline under the CRPS proper score and produces the best-calibrated predictive distributions of any method evaluated, in a literature where every scored precedent is underdispersed. Its training memory is constant from 64^3 to 1024^3 output resolution, where voxel diffusion and the autoencoder hit measured walls at 320^3 and 640^3 , and its inference is constant-memory with time linear in the number of evaluation points. Under noisy, occluded, and channel-dropped sensing on wildfire simulation its worst-case error is below every baseline’s best case, and it infers volumetric flow from surface-only sensors on unstructured transonic-wing meshes that grid-based generative methods cannot represent.

1 Introduction

Volumetric fields govern systems that experiments can only observe pointwise: velocimetry resolves planes of a turbulent flow while the surrounding volume goes unmeasured, and flight-test campaigns record surface pressure at taps while the off-body flow that determines performance is inaccessible. Recovering the full three-dimensional state is a long-standing inverse problem (Callaham et al., 2019; Raissi et al., 2020; Vinuesa et al., 2023), and at realistic sensor densities it is severely ill-posed: many fields fit the data, and the missing information sits disproportionately at the small scales. Deterministic reconstruction networks (Fukami et al., 2021; Santos et al., 2023; Dang et al., 2025) return a conditional mean—the optimal point estimate and a systematically over-smoothed one, blurring precisely the fine structure whose recovery motivates reconstruction. Generative models resolve the mismatch by sampling fields consistent with the observations, retaining sharp realizations and quantifying what the measurements leave undetermined (Shu et al., 2023; Li et al., 2024; Zhuang et al., 2025).

This posterior-sampling program has largely been carried out in two dimensions (Shu et al., 2023; Amorós-Trepat et al., 2026; Li et al., 2024; Huang et al., 2024; Zhuang et al., 2025). In parallel, generative modeling has been lifted from fixed grids to function space, defining flows

over functions that are consistent across discretizations (Kerrigan et al., 2024b; Lim et al., 2025; Pidstrigach et al., 2024; Shi et al., 2025), with guided variants for inverse problems (Yao et al., 2025; Kerrigan et al., 2024a). These methods establish the right abstraction for scattered measurements, but their demonstrations remain at 64×64 to 128×128 on uniform grids with spectral backbones that presuppose them; scaling guided function-space sampling beyond two dimensions is stated as an open problem (Yao et al., 2025).

The transition to three dimensions has been attempted along two detours, each with a documented cost. The first is voxelization: dense 3D U-Nets and transformers compute globally on the full grid at every step, so activation memory grows with the cube of resolution, and published 3D generative models cap their domains at $192 \times 48 \times 48$ (Lienen et al., 2024), train on 32^3 crops with attention removed (Oommen et al., 2026), factorize attention at 64^3 (Li et al., 2023a), or resort to patch decompositions (Holzschuh et al., 2026); it also forfeits the geometry, since unstructured meshes must first be resampled onto axis-aligned grids. The second is latent compression (Du et al., 2024; Zhou et al., 2025; Wang et al., 2026; Li et al., 2025), which makes 3D generation affordable but bounds reconstruction by the autoencoder, and that bound is not tight enough for turbulence: MSE-trained autoencoders retain roughly a quarter of the spectral power in the dissipation range (Rafiq & Nair, 2026; Glaws et al., 2020), and latent neural-field models report small-scale discrepancies on their own benchmarks (Du et al., 2024). A model whose prior cannot represent the small scales cannot return them, whatever the observations. Meanwhile the architectures that do scale to realistic 3D discretizations (Luo et al., 2025; Alkin et al., 2024; 2025; Li et al., 2023b) are uniformly deterministic, and generative modeling on meshes has been limited to a few thousand nodes (Lino et al., 2025).

A further gap separates published evaluations from deployed sensing. The dominant protocol conditions on noiseless values at freely chosen interior locations (Shu et al., 2023; Huang et al., 2024; Amorós-Trepat et al., 2026), but real measurements are noisy, often confined to surfaces (Guastoni et al., 2021; Cuéllar et al., 2024; Mousavi & Eldredge, 2025), occluded by the phenomena being measured, and frequently of different quantities than the ones sought (Barwey et al., 2021; Kildare et al., 2024). That regime is where the generative formulation earns its cost: with noiseless, well-placed observations deterministic models are competitive (Zhuang et al., 2025).

We introduce LOCUS, which performs conditional rectified flow directly on field values at scattered points—ambient space, end to end, no autoencoder and no voxel grid—with a Gaussian-process source that ties the pointwise flow into a coherent flow over functions. Tractability follows from two decisions. First, conditioning is factored by locality: a cross-attention encoder compresses the sensor set into a compact global summary at cost linear in the number of sensors, and each query retrieves a small, importance-weighted neighborhood of sensor tokens at constant cost, so the network never forms a global volumetric representation. Second, the flow-matching objective is a domain integral admitting unbiased estimation on random point subsets, so LOCUS trains on roughly two percent of the discretization per step, decoupling training memory from resolution. Neither decision is available to a voxel U-Net or patchified transformer, whose computation graphs are global in the discretization. Conditioning is amortized over randomized sensor configurations, so one model serves a broad family of counts, placements, noise levels, and observed-channel subsets—including operators beyond the training ranges (Section 4.3)—and rectified flow yields samples in two to sixteen network evaluations.

Our contributions support one central thesis: ambient pointwise conditional generation removes the 3D discretization bottleneck of generative field reconstruction, without latent compression. Specifically:

- **Model.** The first generative model reconstructing 3D fields at realistic discretization sizes directly in ambient function space, combining locality-factored conditioning, subsampled function-space training, and a binned spectral objective on the rectified-flow endpoint (Section 3).
- **Measured scaling separation.** Training memory and step time are constant in output discretization from 64^3 to 1024^3 (39.9 GB, 0.44 s) and inference is constant-memory (1.4 GB) with time linear in point count, while voxel diffusion and the latent base-

line’s autoencoder hit measured out-of-memory walls at 320^3 and 640^3 : grid methods must supervise every cell at every step, whereas LOCUS supervises a chosen 39k-point budget (Section 4.2).

- Accuracy with calibration under honest evaluation. On a spatially disjoint held-out region of forced isotropic turbulence—a protocol adopted after showing the standard random-in-time split reduces to near-duplicate interpolation—LOCUS beats a parameter-matched latent flow-matching baseline on CRPS at one operating point and is TOST-equivalent at the other, concedes at most a small bounded relative- L_2 deficit, and is decisively better calibrated ($t \approx 9$ –28 over 50 paired snapshots); a retrained voxel-diffusion baseline trails on both axes (Section 4.1).
- Deployed sensing. Under realistic operators on wildfire LES (noise, slab occlusion, channel dropout), one amortized LOCUS degrades gracefully—its worst case (0.60) beats the best baseline’s clean case (0.72)—infers unobserved fields through the same interface, and, on 673 transonic wing meshes, infers volumetric predictive distributions from surface-only sensors, an operator grid-based generative baselines cannot represent (Section 4.3). Throughout we treat the sampled distributions as approximations to the posterior whose calibration we measure rather than assume.

2 Related Work

Sparse field reconstruction. Deterministic learned reconstruction spans convolutional models on Voronoi-embedded sensor images (Fukami et al., 2021), attention over arbitrary sensor sets (Santos et al., 2023; Dang et al., 2025), shallow decoders (Erichson et al., 2020), and wall-to-volume estimation (Guastoni et al., 2021; Cuéllar et al., 2024; Mousavi & Eldredge, 2025); these return point estimates that regress toward the conditional mean, and classical assimilation (Raissi et al., 2020; Cai et al., 2021) recovers fields per case without amortization. Generative counterparts condition diffusion through sampling-time guidance (Shu et al., 2023; Huang et al., 2024; Amorós-Trepát et al., 2026; Chung et al., 2023; Rozet & Louppe, 2023), pretrained score models (Li et al., 2024), or amortized training (Zhuang et al., 2025; Liu & Thuerey, 2024), almost exclusively on uniform 2D grids; the 3D exceptions pay the voxel cost (Oommen et al., 2026; Lienen et al., 2024), and latent alternatives (Du et al., 2024; Zhou et al., 2025; Lino et al., 2025) inherit the autoencoder’s spectral floor (Rafiq & Nair, 2026; Glaws et al., 2020).

Probabilistic evaluation of reconstruction. Although this literature routinely draws ensembles, it almost never scores them: evaluation stops at pointwise error and physics statistics, and several works claim calibrated uncertainty on the evidence of standard-deviation maps alone (Wang et al., 2025; Parikh et al., 2026; Yang et al., 2025). Proper scoring rules and calibration diagnostics (Gneiting & Raftery, 2007) appear only at the margins: a latent-diffusion boundary-layer reconstruction reports a rank histogram and finds its ensemble underdispersed (Rybachuk et al., 2023), kilometer-scale generative weather assimilation scores CRPS and spread–skill and likewise finds miscalibration (Manshausen et al., 2025), graph-diffusion urban assimilation reports CRPS on one fixed geometry (Giral et al., 2026), and sensor-time-history decoders report CRPS and coverage for 2D slices (Gao et al., 2026). To our knowledge no prior work demonstrates a calibrated distribution for 3D reconstruction from scattered point sensors under a proper scoring rule; every scored precedent in any adjacent setting is underdispersed. We therefore report fair CRPS, spread–error, coverage, and rank histograms for our model and every baseline that can produce an ensemble, with paired per-snapshot tests.

Function-space generative models and 3D scaling. Diffusion and flow matching have been lifted to function spaces with well-posedness guarantees (Lim et al., 2025; Pidstrigach et al., 2024; Kerrigan et al., 2024b; Franzese et al., 2023; Shi et al., 2024; 2025), with guided variants for inverse problems (Yao et al., 2025; Baldassari et al., 2023; Kerrigan et al., 2024a) and physics instantiations (Wang et al., 2026; Bond-Taylor & Wilcocks, 2024), but demonstrations are one- and two-dimensional with FFT backbones tied to uniform grids. Deterministic surrogates separately reach millions of mesh points through graph, linear-

attention, latent-grid, and anchored-attention designs (Pfaff et al., 2021; Hao et al., 2023; Wu et al., 2024; Li et al., 2023b; Luo et al., 2025; Alkin et al., 2024; 2025), establishing that locality is what scales; none of that line is generative. LOCUS keeps the function-space formulation while transferring those scaling ingredients into the conditioning pathway itself, and generates directly in ambient space on the measurement discretization with no compression stage.

3 Method

3.1 Setup: rectified flow on point sets with a Gaussian-process source

Let $u : \Omega \rightarrow \mathbb{R}^C$ be a physical field on $\Omega \subset \mathbb{R}^3$, drawn from a distribution over fields (turbulent-flow snapshots, RANS solutions over varying geometry). It is observed through a measurement operator returning readings $\mathcal{O} = \{(\mathbf{x}_j, c_j, y_j)\}_{j=1}^M$ with $y_j = u_{c_j}(\mathbf{x}_j) + \varepsilon_j$, where $\mathbf{x}_j$ is a sensor location, c_j indexes the observed quantity, and ε_j is noise. This covers locations scattered arbitrarily or confined to a surface, observed quantities forming a subset of the target channels or different quantities altogether, and known covariates entering as additional tokens. The goal is to sample an approximation to $p(u \mid \mathcal{O})$ at arbitrary query locations $\{\mathbf{y}_i\}_{i=1}^N$ that need not coincide with any training discretization; how well that approximation holds is an empirical question we answer with calibration diagnostics rather than assume.

LOCUS models the conditional distribution with a rectified flow (Liu et al., 2023; Lipman et al., 2023; Albergo & Vanden-Eijnden, 2023) defined directly on field values. The source distribution is a Gaussian process rather than pointwise white noise: a draw $u_0 \sim \pi_0$ evaluated at coordinates $\{\mathbf{y}_i\}$ is produced by a random-feature expansion (Rahimi & Recht, 2007),

$$u_0(\mathbf{y}) = \sqrt{\frac{2}{F}} \sum_{f=1}^F w_f \cos(\boldsymbol{\omega}_f^\top \mathbf{y} + b_f), \quad w_f \sim \mathcal{N}(0, 1), \quad (1)$$

with frequencies $\boldsymbol{\omega}_f$ drawn once from a Gaussian whose scale sets the prior correlation length. Evaluating equation 1 at any two point sets yields marginals of one consistent function-valued source, which makes training and sampling coherent across discretizations. Given a data field u_1 at the same points, the state interpolates linearly, $u_t = (1-t)u_0 + tu_1$, and a velocity network v_θ regresses the constant displacement:

$$\mathcal{L}_{\text{FM}}(\theta) = \mathbb{E}_{t, u_0, u_1, \mathcal{O}} \frac{1}{|\Omega|} \int_{\Omega} \|v_\theta(t, u_t(\mathbf{y}), \mathbf{y}; \mathcal{O}) - (u_1(\mathbf{y}) - u_0(\mathbf{y}))\|_2^2 d\mathbf{y}. \quad (2)$$

The objective is a domain integral, and this removes the resolution barrier: equation 2 admits an unbiased Monte Carlo estimate over a random subset of query points at every step. LOCUS trains on roughly 2% of the discretization per step (3.9×10^4 of 1.95×10^6 points on the turbulence benchmark), so training memory is set by the subset size, not the resolution; voxel U-Nets, latent autoencoders, and patchified transformers do not admit this estimator, because their graphs couple all grid locations. Sampling integrates the learned ODE from a fresh source draw with a uniform Euler scheme; rectified-flow paths are near-straight, so 2–16 evaluations suffice against the hundreds used by guided diffusion (Li et al., 2024; Huang et al., 2024), and repeated integration under the same $\mathcal{O}$ yields an ensemble.

3.2 Locality-factored conditioning

The velocity network must expose the observations to every query point without incurring an all-to-all term in N or a dense volumetric representation. LOCUS factors the conditioning pathway by locality (Figure 1).

Each observation $(\mathbf{x}_j, c_j, y_j)$ is embedded as a token from a Fourier positional encoding of $\mathbf{x}_j$ (Tancik et al., 2020), the reading y_j , and a learned embedding of the quantity index c_j . A small array of S learnable latents cross-attends to the M sensor tokens (Jaegle et al., 2021) at cost $\mathcal{O}(SM)$, refined by self-attention blocks that re-attend to the sensor tokens




Figure 1: The LOCUS architecture. Sensor tokens are compressed by cross-attention into a small latent array (global pathway, $\mathcal{O}(SM)$), which re-injects context back into the sensor tokens. Each query assembles its velocity from its point features, a readout from the latent array, and an importance-weighted retrieval over its K nearest sensor tokens (local pathway, $\mathcal{O}(NK)$). No dense volumetric representation is formed; cost is linear in the number of queries, and training estimates the loss on a random $\sim 2\%$ point subset per step.

between rounds; the mean-pooled array provides a global summary, and a final pass writes that context back into the sensor tokens.

Each query y_i then fuses three components: a point feature encoding its position, the current state $u_t(y_i)$, and t ; a global readout cross-attending from a coordinate-conditioned decoder token to the latent array at cost $\mathcal{O}(L)$; and a local retrieval over the K nearest sensor tokens under an importance-warped metric $d(y_i, x_j)^2 = \|y_i - x_j\|_2^2 - \beta_j$, combined with normalized RBF weights. The learned per-sensor bias β_j lets informative sensors extend their influence beyond their Euclidean neighborhood, which matters when sensing is anisotropic (surface-only layouts) or partially corrupted. A FiLM-modulated coarse prediction from the global summary enters as a gated residual, and a small MLP head produces the velocity. Retrieval and the k -nearest-neighbor search run as streaming kernel reductions (Charlier et al., 2021), so the $[N \times M]$ distance matrix is never materialized. Total cost per evaluation is $\mathcal{O}(SM + NL + NK)$ with no term coupling query points to each other: compute is linear in N , and the same trained network evaluates on 10^4 or 2×10^6 points unchanged.

3.3 Training objective and sampling cost

Two additions to equation 2 matter in practice. First, t is drawn logit-normally, concentrating capacity at intermediate noise levels (Ma et al., 2024). Second, generative models of turbulence exhibit spectral bias, reproducing large scales and attenuating the inertial range (Rahaman et al., 2019); following the observation that binned spectral penalties succeed where naive FFT losses do not (Chakraborty et al., 2026), each batch appends a 32^3 lattice block to the query set, and from the endpoint estimate $\hat{u}_1 = u_t + (1 - t)v_\theta$ we penalize the squared difference of shell-averaged log power against u_1 over 12 radial bins, weighted by t . The total loss is $\mathcal{L}_{\text{FM}} + \lambda_{\text{sp}}\mathcal{L}_{\text{sp}}$ with $\lambda_{\text{sp}} = 0.02$, and Section 4.1 reports the resulting trade-off with λ_{sp} as the control.

At sampling time every quantity depending only on the observations and query coordinates—the sensor branch, the neighbor geometry, and the coordinate-conditioned readout—is computed at the first ODE step and reused exactly thereafter, so the marginal cost of a step falls to the fusion head: a full 125^3 sample costs 4.5s at 4 evaluations and 5.1s at 16, making high-step operating points nearly free.

Amortized measurement operators. Conditioning is amortized by randomizing the measurement operator during training: each step draws the sensor count and layout afresh, and with configurable probability readings are corrupted with Gaussian noise of random amplitude, contiguous sensor regions are removed, and entire observed channels are dropped. The network thus learns one conditional velocity field over a family of operators, and at test time a single model serves that family without retraining or guidance; Section 4.3 additionally probes operators outside the training ranges. Known covariates (Mach, incidence) enter as parameter tokens exempt from corruption. Observed values can be enforced on the generated field at sampling time; we enforce for noiseless protocols and disable under sensor noise, where interpolating corrupted readings is undesirable. Training details are in Appendix B.

4 Experiments

Evaluation protocol. Turbulence snapshots are strongly autocorrelated in time, and splitting consecutive DNS frames at random—standard practice—makes every validation frame a near-duplicate of a training frame (frame-to-frame correlation 1.00 at unit separation). Under such a split all methods, ours included, achieve errors two to three times lower, and the ranking measures interpolation rather than reconstruction. All turbulence results therefore use spatial holdout: training on three 125^3 regions of the DNS, evaluation on a fourth, disjoint region, with temporally blocked splits inside the training set. Probabilistic metrics use $K=8$ -member ensembles on matched test snapshots with identical seeded sensor draws for every method; comparisons use paired per-snapshot confidence intervals. A temporally blocked same-region protocol accompanies the spatial holdout so that both the conventional setting and the harder extrapolative one are represented. Exact octahedral symmetry augmentation (48 signed axis permutations, discretization-exact on cubic grids) is supplied to every method, ours and baselines alike.

Datasets. Three domains, each exposing one failure mode of existing methods (preprocessing in Appendix C). (i) Forced isotropic turbulence from the Johns Hopkins Turbulence Database (Li et al., 2008): four disjoint 125^3 cutouts of the 1024^3 DNS (1.95×10^6 points each; velocity and pressure; 50 coarsely strided snapshots per cutout), three for training and the fourth for evaluation. (ii) SHIFT-WING (Luminary Cloud, 2026): adaptively meshed RANS solutions over parametric transport-aircraft wing-fuselage variants at Mach 0.5–0.85; 673 cases (600/73), each with 4×10^5 volume nodes and 6.5×10^4 surface nodes carrying pressure and wall shear stress, with geometry varying across cases and entering only through the observation point cloud. (iii) FireBench (Wang et al., 2024): wildfire large-eddy simulations at two wind speeds; we extract the fire-front region (3.7×10^6 points; velocity, potential temperature, fuel density; 120 snapshots).

Baselines. We compare against the routes identified in Section 1, configured for fairness rather than convenience (Appendix D). Latent route: a 3D convolutional autoencoder with rectified flow in its latent space (LDM-style; Rombach et al., 2022), sensor-conditioned through a point encoder, parameter-matched to LOCUS, trained in two stages on identical data, splits, and augmentation; it stands in for the latent-diffusion family (Du et al., 2024). Voxel route: Gen4Turb (Oommen et al., 2026), a published diffusion U-Net with released code, run at its native 120^3 on the same data, wall-clock-matched, and additionally retrained under our exact observation protocol. Deterministic point route: Senseiver (Santos et al., 2023), trained with identical randomization and augmentation.

Metrics. We report relative L_2 of the ensemble mean and of single samples (the latter measures whether individual realizations are physical rather than over-smoothed), band-resolved spectral ratios, fair CRPS (Gneiting & Raftery, 2007), spread-error ratio, and central 90% coverage, all in per-field standardized units; definitions are in Appendix E. Calibration estimates from $K=8$ ensembles are coarse; the fair-CRPS estimator corrects the leading ensemble-size bias, and a $K=32$ check shows the $K=8$ coverage estimates are conservative for every method (ours $0.55 \rightarrow 0.66$, latent baseline $0.45 \rightarrow 0.52$ on the same snapshots) with the ranking and spread-error ratios unchanged.

4.1 Accuracy and calibration on spatially disjoint turbulence

Table 1 reports the head-to-head, with paired per-snapshot statistics over all 50 test snapshots. On the proper score LOCUS wins or ties: without the spectral penalty its CRPS beats the latent baseline significantly (-0.014 , 95% CI $[-0.026, -0.002]$), and the primary $\lambda_{sp}=0.02$ model is indistinguishable (-0.011 , CI $[-0.025, +0.004]$), with TOST equivalence inside a ± 0.03 margin ($p=0.004$). On ensemble-mean relative L_2 the latent baseline holds a small but significant edge ($+0.029$, CI $[+0.009, +0.050]$), bounded below 0.05 by TOST ($p=0.02$). Calibration separates decisively: spread-error and coverage differences of $+0.097$ $[+0.075, +0.118]$ and $+0.087$ $[+0.069, +0.106]$ ($t \approx 9$; for $\lambda_{sp}=0$, $+0.29$ and $+0.20$ at $t \approx 28$). No model here is calibrated in the absolute sense—coverage stays below nominal, as in every

Table 1: Isotropic turbulence, spatially disjoint evaluation region (1.95M points, two velocity channels observed at 1% of points, $K=8$, all 50 held-out snapshots). Spread–error and coverage ideals are 1.0 and 0.90. Senseiver is deterministic (CRPS reduces to MAE, no spread). Gen4Turb is retrained under our observation protocol at its native 120^3 ; under its native shared-mask protocol, which observes all four channels (twice the information), it reaches 0.354/0.204 on an 8-snapshot subset.

Method	Rel. L_2 mean	Rel. L_2 sample	CRPS	Spread/Err.	Cov. _{.90}
Senseiver (deterministic)	0.722	—	0.542	—	—
Gen4Turb (voxel diffusion)	0.713	0.735	0.441	0.28	0.25
Latent flow matching	0.564	0.629	0.302	0.53	0.45
LOCUS ($\lambda_{sp}=0$, NFE 4)	0.623	0.811	0.288	0.82	0.65
LOCUS (NFE 2)	0.582	0.641	0.312	0.37	0.33
LOCUS (NFE 4)	0.593	0.711	0.291	0.63	0.54




Figure 2: Turbulence statistics of single posterior samples on the held-out region (4 snapshots, NFE 4; ensemble means are deliberately not shown—averaging destroys small scales by construction). Left: shell-averaged energy spectra; the latent baseline carries more inertial-range energy, but truncates in the dissipation band (shaded), where LOCUS tracks the DNS. Middle: longitudinal second-order structure functions. Right: velocity-gradient PDFs, where the latent baseline distorts the shoulders while both LOCUS variants track the DNS shape.

scored precedent, which is why we call these predictive rather than posterior distributions—but LOCUS is closest by a wide margin (0.82 spread–error, 0.65 coverage at $\lambda_{sp}=0$, against 0.53 and 0.45). The voxel-diffusion baseline does not close the gap even retrained under the matched protocol: 0.713 relative L_2 with a severely underdispersed ensemble (coverage 0.25), and retraining moved accuracy by less than 0.01—a mask-channel voxel U-Net extracts little of the information in scattered point observations. Its stronger shared-mask number (0.354) requires observing all four channels, which our protocol deliberately does not grant.

Two knobs expose a distortion–dispersion frontier. Step count: two evaluations give the sharpest single samples (0.641, best of any generative method), four are the CRPS optimum, and dispersion rises monotonically with further steps (spread–error 0.92 at NFE 16 for $\lambda_{sp}=0$, at flat CRPS; Appendix F), which the per-solve conditioning cache makes nearly free. Spectral weight: raising λ_{sp} from 0 to 0.02 improves single-sample error and inertial-band fidelity at a measured cost in spread, and at 0.05 the accuracy gain saturates while dispersion degrades further. We report $\lambda_{sp}=0.02$ as the primary model; no single point on the frontier is best for every use.

Two potential confounds are closed by ablation (Appendix F). Doubling the per-step query budget from 2% to 4% of the field does not improve held-out loss (three seeds), so subsampled supervision is not the binding constraint, even though the latent baseline consumes roughly $50\times$ more supervised values per step by decoding the full grid; and the octahedral augmentation supplied to all methods is the operative one, with continuous rotations, translations, and proper-rotation restrictions all falling within run-to-run variance.

Small-scale fidelity. Figure 2 tests the premise behind avoiding an autoencoder and returns a more specific answer than the premise assumed. In band-integrated energy relative to the




Figure 3: Measured cost versus output resolution (H100, bf16, log-log; OOM markers are measured failures, not extrapolations), for training steps (top) and full-field posterior sampling (bottom). Grid methods (voxel diffusion, convolutional AE) scale cubically and fail at 320^3 and 640^3 ; point methods (LOCUS, Senseiver) are flat in memory to 1024^3 (10^9 points). Annotations give the fraction of the field supervised per training step: 100% forced for grid methods, a fixed 39k-point budget for LOCUS (2% at 125^3 , 0.004% at 1024^3).

DNS, the latent baseline retains more inertial-range energy than LOCUS (0.73 vs. 0.56 over $k \in [8, 31]$) but truncates the dissipation range (0.81 vs. 0.97 over $k \in [32, 62]$): compression costs the smallest scales specifically, where the autoencoder’s spectral floor was predicted to bite (Rafiq & Nair, 2026; Glaws et al., 2020), while the ambient model’s deficit sits at intermediate scales. The spectral objective moves LOCUS in the intended direction on both bands ($0.47 \rightarrow 0.56$, $0.88 \rightarrow 0.97$), and the velocity-gradient PDFs show the latent baseline distorting the shoulders of the distribution while both LOCUS variants track the DNS shape. This is a crossover, not a uniform win: neither model reproduces the full spectrum from 1% sensors, and what we can claim is that ambient generation preserves dissipation-range content that latent compression removes.

4.2 The cost of resolution

Figure 3 reports measured peak memory and wall-clock for training steps and full-field sampling from 64^3 to 1024^3 ; every out-of-memory marker is an observed failure on an 80 GB H100. The axis separates two computational families. Grid methods scale cubically: the convolutional autoencoder anchoring the latent baseline trains at 50.9 GB for one 512^3 field and fails at 640^3 at batch size 1 (at production batch size the wall falls between 256^3 and 320^3), and the voxel diffusion model fails at 320^3 at batch size 1 while already consuming 57.5 GB at 120^3 . Point methods are flat: LOCUS trains at 39.9 GB and 0.44s per step and samples at 1.4 GB peak at every resolution tested, and Senseiver shares the profile (19.4 GB train, 1.2 GB inference)—flatness is a property of point computation, not of our model alone. Within that family LOCUS is the only generative member, and the deterministic one pays a factor of 1.3 in accuracy and cannot produce an ensemble.

The annotations give the mechanism: a voxel or convolutional model is structurally committed to supervising every grid cell at every step—its memory is the supervision—whereas LOCUS chooses its budget independently of resolution, and the ablations above show that budget saturates. The honest converse is that grid computation is cheaper below the crossover ($\sim 240^3$ for inference): a convolutional decode of a 125^3 field costs 6 ms where our chunked point evaluation costs 4.5 s. Stated precisely, the claim is constant memory with respect to the output discretization and linear-time inference, not constant total cost; the regime where LOCUS is the only option is the one the applications occupy—resolutions above $\sim 300^3$, unstructured discretizations, and sensing geometries no grid can represent. The constancy of the training graph also raises the right follow-up question, whether the model trained at

Table 2: Wildfire LES (3.7M points, 5 fields) under realistic measurement operators (relative L_2 of ensemble mean / fair CRPS; $K=8$, 8 snapshots, identical seeded operator realizations for all methods). Sensors observe the three wind components at $\sim 1\%$ of points. LOCUS is a single amortized model across all rows.

Operator	LOCUS		Latent FM		Senseiver	
	Rel. L_2	CRPS	Rel. L_2	CRPS	Rel. L_2	CRPS
Clean	0.456	0.234	0.724	0.356	0.779	0.443
Noise $\sigma=0.1$	0.458	0.237	0.722	0.355	0.779	0.444
Noise $\sigma=0.3$	0.461	0.248	0.724	0.356	0.779	0.445
Slab occlusion (25%)	0.511	0.275	0.720	0.355	0.780	0.443
Channel dropout (u only)	0.603	0.355	0.719	0.309	0.908	0.514

one discretization transfers as a function-valued sampler to finer and irregular ones; a direct discretization-transfer experiment (training at 125^3 , reconstructing a native 500^3 cutout and irregular subsets of it) is in preparation as a core test of the function-space claim.

4.3 Deployed sensing: corrupted operators and native geometry

Table 2 evaluates the regime the introduction argued for, on wildfire LES at 3.7M points. One amortized LOCUS wins every cell, with qualitative margins: 0.3σ sensor noise costs one percent relative, removing a 25% slab of sensors twelve percent, and dropping two of three wind components a third. Its worst case (0.603) still beats either baseline’s best case under clean, complete observations (0.719 and 0.779). One row is genuinely out of distribution—training used noise amplitudes up to 0.1σ , so the 0.3σ row is triple anything seen—while the occlusion and dropout rows lie inside the amortized family, demonstrating robustness within a deliberately broad operator distribution rather than extrapolation.

The baselines’ apparent stability across rows deserves the correct reading: their numbers barely move because they are bias-dominated, and a model that extracts little of the information in the sensors loses little when that information is corrupted. The latent baseline reconstructs at 0.72 relative L_2 under every operator (predicting the zero field scores 1.0), and its CRPS improves under channel dropout only because its spread is large where its mean is wrong. LOCUS’s errors respond to the information content of the observations because the model uses that information; graceful, monotone degradation, not invariance, is the behavior a practitioner should want. The latent baseline’s dispersion ratio here (0.87) exceeds ours (0.76), but at more than twice the CRPS; the proper score, integrating both, favors LOCUS by 34% in every condition. A clean-trained control isolates what operator amortization actually buys: trained without any corruption, the control matches the robust model under noise and slab occlusion (0.458 vs. 0.461 at 0.3σ ; 0.513 vs. 0.511 under occlusion)—robustness to those perturbations comes from the conditioning architecture, not the training distribution—but collapses under channel dropout (0.743 vs. 0.603), the one operator that changes the structure of the observation set. Amortized operator training is what covers that structural shift, and it costs nothing in the clean condition (0.455 vs. 0.456).

Cross-variable inference falls out of the same interface: observing only wind, the model recovers unobserved fuel density at 0.044 relative L_2 and potential temperature at 0.49, degrading only to 0.064 and 0.52 when two of the three wind components are also removed (Appendix F). The randomized observed-channel embedding is what makes a wind-only, single-component operator an interpolation rather than an extrapolation for the conditioning network.

Surface-only observation on native geometry. The wing domain exercises two capabilities no grid-based generative method possesses: operation on an unstructured mesh whose geometry changes across cases, and conditioning on observations confined to the body surface. From 640 surface sensors (512 pressure taps, 128 shear gauges, with Mach and incidence as parameter tokens), LOCUS recovers the volumetric pressure-coefficient field on 8 held-out

geometries at 0.126 relative L_2 and velocities at 0.53–0.69, with CRPS 0.18, spread–error 0.68, and 90% coverage 0.38 ($K=4$, NFE 4); the predictive standard deviation concentrates in the boundary layer and wake, the regions genuinely undetermined by surface data. The importance-warped retrieval is designed to let surface sensors inform off-body queries—under a plain Euclidean metric most volume points have no nearby sensors—and we treat its necessity as a hypothesis for direct ablation. Grid-based baselines are structurally excluded here, which is the point.

5 Conclusion

Generative reconstruction of 3D fields has been blocked not by modeling capacity but by a structural mismatch: generative architectures compute globally on dense discretizations, while the problem is defined on scattered points and needs only local and summary information at each. LOCUS resolves it by carrying rectified flow directly on field values with a Gaussian-process source, locality-factored conditioning, and subsampled function-space training. The measured consequences: costs flat from 64^3 to 1024^3 where voxel and autoencoder pipelines fail at 320^3 – 640^3 ; accuracy competitive with a strong latent baseline on spatially disjoint turbulence, with the best-calibrated predictive distributions of any method evaluated; uniform dominance under corrupted sensing at 3.7M points; and volumetric predictions from surface-only sensors on varying unstructured geometries.

Limitations are equally measurable. Coverage stays below nominal even at our best-calibrated operating point (0.65 against 0.90)—best-calibrated among evaluated methods, not calibrated; single-snapshot reconstruction ignores temporal information that assimilation exploits; per-domain training is still required; and constant-memory inference pays a linear time cost, making LOCUS the slowest method in absolute terms at large resolution (41 minutes for a 10^9 -point sample). Sharpening calibration, conditioning on sensor time histories, and pretraining across flow families are the natural continuations.

Reproducibility statement

All experiments use publicly available data: the JHTDB isotropic turbulence database, the SHIFT-WING dataset (CC-BY-NC-4.0), and the public FireBench store; exact cutout indices, crop definitions, subsampling seeds, and normalization are specified in Appendix C. Architecture and training hyperparameters for LOCUS and all baselines are listed in Appendices A–D together with the configuration files used for every run. Code, configurations, and dataset preparation scripts will be released.

AI use statement

References

- Michael S Albergo and Eric Vanden-Eijnden. Building normalizing flows with stochastic interpolants. In *International Conference on Learning Representations*, 2023.
- Benedikt Alkin, Andreas Fürst, Simon Schmid, Lukas Gruber, Markus Fellner, and Johannes Brandstetter. Universal physics transformers: A framework for efficiently scaling neural operators. In *Advances in Neural Information Processing Systems*, volume 37, 2024.
- Benedikt Alkin, Maurits Bleeker, Richard Kurle, Tobias Kronlachner, Reinhard Sonnleitner, Matthias Dorfer, and Johannes Brandstetter. AB-UPT: Scaling neural CFD surrogates for high-fidelity automotive aerodynamics simulations via anchored-branched universal physics transformers. *Transactions on Machine Learning Research*, 2025.
- Marc Amorós-Trepat, Luis Medrano-Navarro, Qiang Liu, Luca Guastoni, and Nils Thuerey. Guiding diffusion models to reconstruct flow fields from sparse data. *Physics of Fluids*, 38(1):015112, 2026.
- Lorenzo Baldassari, Ali Siahkoohi, Josselin Garnier, Knut Solna, and Maarten V de Hoop. Conditional score-based diffusion models for Bayesian inference in infinite dimensions. In *Advances in Neural Information Processing Systems*, volume 36, 2023.

- Shivam Barwey, Venkat Raman, and Adam M Steinberg. Extracting information overlap in simultaneous OH-PLIF and PIV fields with neural networks. *Proceedings of the Combustion Institute*, 38(4):6241–6249, 2021.
- Sam Bond-Taylor and Chris G Willcocks. ∞ -Diff: Infinite resolution diffusion with subsampled mollified states. In *International Conference on Learning Representations*, 2024.
- Shengze Cai, Zhicheng Wang, Frederik Fuest, Young Jin Jeon, Callum Gray, and George Em Karniadakis. Flow over an espresso cup: inferring 3-D velocity and pressure fields from tomographic background oriented Schlieren via physics-informed neural networks. *Journal of Fluid Mechanics*, 915:A102, 2021.
- Jared L Callahan, Kazuki Maeda, and Steven L Brunton. Robust flow reconstruction from limited measurements via sparse representation. *Physical Review Fluids*, 4(10):103907, 2019.
- Dibyajyoti Chakraborty, Arvind T. Mohan, and Romit Maulik. Binned spectral power loss for improved prediction of chaotic systems. *Journal of Computational Physics*, 2026. arXiv:2502.00472.
- Benjamin Charlier, Jean Feydy, Joan Alexis Glaunès, François-David Collin, and Ghislain Durif. Kernel operations on the GPU, with autodiff, without memory overflows. *Journal of Machine Learning Research*, 22(74):1–6, 2021.
- Hyungjin Chung, Jeongsol Kim, Michael T Mccann, Marc L Klasky, and Jong Chul Ye. Diffusion posterior sampling for general noisy inverse problems. In *International Conference on Learning Representations*, 2023.
- Antonio Cuéllar, Alejandro Güemes, Andrea Ianaro, Óscar Flores, Ricardo Vinuesa, and Stefano Discetti. Three-dimensional generative adversarial networks for turbulent flow estimation from wall measurements. *Journal of Fluid Mechanics*, 991:A1, 2024.
- Hiep Vo Dang, Jun Won Choi, and Phong CH Nguyen. FLRONet: deep operator learning for high-fidelity fluid flow field reconstruction from sparse sensor measurements. *Journal of Fluid Mechanics*, 1008:A25, 2025.
- Pan Du, Meet Hemant Parikh, Xiantao Fan, Xin-Yang Liu, and Jian-Xun Wang. Conditional neural field latent diffusion model for generating spatiotemporal turbulence. *Nature Communications*, 15:10416, 2024.
- N Benjamin Erichson, Lionel Mathelin, Zhewei Yao, Steven L Brunton, Michael W Mahoney, and J Nathan Kutz. Shallow neural networks for fluid flow reconstruction with limited sensors. *Proceedings of the Royal Society A*, 476(2238):20200097, 2020.
- Giulio Franzese, Giulio Corallo, Simone Rossi, Markus Heinonen, Maurizio Filippone, and Pietro Michiardi. Continuous-time functional diffusion processes. In *Advances in Neural Information Processing Systems*, volume 36, 2023.
- Kai Fukami, Romit Maulik, Nesar Ramachandra, Koji Fukagata, and Kunihiko Taira. Global field reconstruction from sparse sensors with Voronoi tessellation-assisted deep learning. *Nature Machine Intelligence*, 3:945–951, 2021.
- Mars Liyao Gao, Yuxuan Bao, Amy S Rude, Xinwei Shen, and J Nathan Kutz. UQ-SHRED: uncertainty quantification of shallow recurrent decoder networks for sparse sensing via engression. *arXiv preprint arXiv:2604.01305*, 2026.
- Francisco Giral, Alberto Manzano, Pablo Gómez, Ricardo Vinuesa, and Soledad Le Clainche. Generative data assimilation on complex urban areas via classifier-free diffusion guidance. *arXiv preprint arXiv:2601.11440*, 2026.
- Andrew Glaws, Ryan King, and Michael Sprague. Deep learning for in situ data compression of large turbulent flow simulations. *Physical Review Fluids*, 5(11):114602, 2020.

- Tilman Gneiting and Adrian E Raftery. Strictly proper scoring rules, prediction, and estimation. *Journal of the American Statistical Association*, 102(477):359–378, 2007.
- Luca Guastoni, Alejandro Güemes, Andrea Ianaro, Stefano Discetti, Philipp Schlatter, Hossein Azizpour, and Ricardo Vinuesa. Convolutional-network models to predict wall-bounded turbulence from wall quantities. *Journal of Fluid Mechanics*, 928:A27, 2021.
- Zhongkai Hao, Zhengyi Wang, Hang Su, Chengyang Ying, Yinpeng Dong, Songming Liu, Ze Cheng, Jian Song, and Jun Zhu. GNOT: A general neural operator transformer for operator learning. In *International Conference on Machine Learning*, pp. 12556–12569, 2023.
- Benjamin Holzschuh, Georg Kohl, Florian Redinger, and Nils Thuerey. P3D: Scalable neural surrogates for high-resolution 3D physics simulations with global context. In *International Conference on Learning Representations*, 2026.
- Jiahe Huang, Guandao Yang, Zichen Wang, and Jeong Joon Park. DiffusionPDE: Generative PDE-solving under partial observation. In *Advances in Neural Information Processing Systems*, volume 37, 2024.
- Andrew Jaegle, Felix Gimeno, Andy Brock, Oriol Vinyals, Andrew Zisserman, and Joao Carreira. Perceiver: General perception with iterative attention. In *International Conference on Machine Learning*, pp. 4651–4664, 2021.
- Gavin Kerrigan, Giosue Migliorini, and Padhraic Smyth. Dynamic conditional optimal transport through simulation-free flows. In *Advances in Neural Information Processing Systems*, volume 37, 2024a.
- Gavin Kerrigan, Giosue Migliorini, and Padhraic Smyth. Functional flow matching. In *International Conference on Artificial Intelligence and Statistics*, pp. 3934–3942, 2024b.
- James AC Kildare, Jiajun Zhu, Tao Yu, Can Ruan, Xue-Song Bai, Zhongshan Li, and Marcus Aldén. Unsupervised deep learning method for temperature reconstruction of turbulent flames from OH and CH₂O planar laser-induced fluorescence. *Proceedings of the Combustion Institute*, 40:105330, 2024.
- Yi Li, Eric Perlman, Minping Wan, Yunke Yang, Charles Meneveau, Randal Burns, Shiyi Chen, Alexander Szalay, and Gregory Eyink. A public turbulence database cluster and applications to study Lagrangian evolution of velocity increments in turbulence. *Journal of Turbulence*, 9:N31, 2008.
- Zeyu Li, Wang Han, Yue Zhang, Qingfei Fu, Jingxuan Li, Lizi Qin, Ruoyu Dong, Hao Sun, Yue Deng, and Lijun Yang. Learning spatiotemporal dynamics with a pretrained generative model. *Nature Machine Intelligence*, 6:1566–1579, 2024.
- Zijie Li, Dule Shu, and Amir Barati Farimani. Scalable transformer for PDE surrogate modeling. In *Advances in Neural Information Processing Systems*, volume 36, 2023a.
- Zijie Li, Saurabh Patil, Francis Ogoke, Dule Shu, Wilson Zhen, Michael Schneier, John R Buchanan Jr, and Amir Barati Farimani. Latent neural PDE solver: A reduced-order modeling framework for partial differential equations. *Journal of Computational Physics*, 524:113705, 2025.
- Zongyi Li, Nikola B Kovachki, Chris Choy, Boyi Li, Jean Kossaifi, Shourya P Otta, Mohammad A Nabian, Maximilian Stadler, Christian Hundt, Kamyar Azizzadenesheli, and Anima Anandkumar. Geometry-informed neural operator for large-scale 3D PDEs. In *Advances in Neural Information Processing Systems*, volume 36, 2023b.
- Marten Lienen, David Lüdke, Jan Hansen-Palmus, and Stephan Günnemann. From zero to turbulence: Generative modeling for 3D flow simulation. In *International Conference on Learning Representations*, 2024.

- Jae Hyun Lim, Nikola B Kovachki, Ricardo Baptista, Christopher Beckham, Kamyar Azizzadenesheli, Jean Kossaifi, Vikram Voleti, Jiaming Song, Karsten Kreis, Jan Kautz, Christopher Pal, Arash Vahdat, and Anima Anandkumar. Score-based diffusion models in function space. *Journal of Machine Learning Research*, 26(158):1–117, 2025.
- Mario Lino, Tobias Pfaff, and Nils Thuerey. Diffusion graph nets: Learning distributions of unsteady flow states on meshes. In *International Conference on Learning Representations*, 2025.
- Yaron Lipman, Ricky TQ Chen, Heli Ben-Hamu, Maximilian Nickel, and Matt Le. Flow matching for generative modeling. In *International Conference on Learning Representations*, 2023.
- Qiang Liu and Nils Thuerey. Uncertainty-aware surrogate models for airfoil flow simulations with denoising diffusion probabilistic models. *AIAA Journal*, 62(8):2912–2933, 2024.
- Xingchao Liu, Chengyue Gong, and Qiang Liu. Flow straight and fast: Learning to generate and transfer data with rectified flow. In *International Conference on Learning Representations*, 2023.
- Luminary Cloud. SHIFT-WING: An open large-scale simulation dataset for transonic wing aerodynamics. 2026. <https://huggingface.co/datasets/luminary-shift/WING>.
- Huakun Luo, Haixu Wu, Hang Zhou, Lanxiang Xing, Yichen Di, Jianmin Wang, and Mingsheng Long. Transolver++: An accurate neural solver for PDEs on million-scale geometries. In *International Conference on Machine Learning*, 2025.
- Nanye Ma, Mark Goldstein, Michael S Albergo, Nicholas M Boffi, Eric Vanden-Eijnden, and Saining Xie. SiT: Exploring flow and diffusion-based generative models with scalable interpolant transformers. In *European Conference on Computer Vision*, 2024.
- Peter Manshausen, Yair Cohen, Peter Harrington, Jaideep Pathak, Mike Pritchard, Piyush Garg, et al. Generative data assimilation of sparse weather station observations at kilometer scales. *Journal of Advances in Modeling Earth Systems*, 17(10), 2025. arXiv:2406.16947.
- Hanieh Mousavi and Jeff D Eldredge. Low-order flow reconstruction and uncertainty quantification in disturbed aerodynamics using sparse pressure measurements. *Journal of Fluid Mechanics*, 1013:A41, 2025.
- Vivek Oommen, Siavash Khodakarami, Aniruddha Bora, Zhicheng Wang, and George Em Karniadakis. Learning turbulent flows with generative models for super-resolution and sparse flow reconstruction. *Nature Communications*, 17:3707, 2026.
- Meet Hemant Parikh, Xiantao Fan, and Jian-Xun Wang. Conditional flow matching for generative modelling of near-wall turbulence with quantified uncertainty. *Journal of Fluid Mechanics*, 1029, 2026. arXiv:2504.14485.
- Tobias Pfaff, Meire Fortunato, Alvaro Sanchez-Gonzalez, and Peter W Battaglia. Learning mesh-based simulation with graph networks. In *International Conference on Learning Representations*, 2021.
- Jakiw Pidstrigach, Youssef Marzouk, Sebastian Reich, and Sven Wang. Infinite-dimensional diffusion models. *Journal of Machine Learning Research*, 25(414):1–52, 2024.
- Danish Rafiq and Aditya G Nair. Spectrally regularized latent flow matching for turbulence generation. arXiv preprint arXiv:2606.11691, 2026.
- Nasim Rahaman, Aristide Baratin, Devansh Arpit, Felix Draxler, Min Lin, Fred Hamprecht, Yoshua Bengio, and Aaron Courville. On the spectral bias of neural networks. In *International Conference on Machine Learning*, pp. 5301–5310, 2019.
- Ali Rahimi and Benjamin Recht. Random features for large-scale kernel machines. In *Advances in Neural Information Processing Systems*, volume 20, 2007.

- Maziar Raissi, Alireza Yazdani, and George Em Karniadakis. Hidden fluid mechanics: Learning velocity and pressure fields from flow visualizations. *Science*, 367(6481):1026–1030, 2020.
- Robin Rombach, Andreas Blattmann, Dominik Lorenz, Patrick Esser, and Björn Ommer. High-resolution image synthesis with latent diffusion models. In *IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pp. 10684–10695, 2022.
- François Rozet and Gilles Louppe. Score-based data assimilation. In *Advances in Neural Information Processing Systems*, volume 36, 2023.
- Alex Rybchuk, Malik Hassanaly, Nicholas Hamilton, Paula Doubrawa, Mitchell J Fulton, and Luis A Martínez-Tossas. Ensemble flow reconstruction in the atmospheric boundary layer from spatially limited measurements through latent diffusion models. *Physics of Fluids*, 35(12):126604, 2023.
- Javier E Santos, Zachary R Fox, Arvind Mohan, Daniel O’Malley, Hari Viswanathan, and Nicholas Lubbers. Development of the Senseiver for efficient field reconstruction from sparse observations. *Nature Machine Intelligence*, 5:1317–1325, 2023.
- Yaozhong Shi, Angela F Gao, Zachary E Ross, and Kamyar Azizzadenesheli. Universal functional regression with neural operator flows. *Transactions on Machine Learning Research*, 2024.
- Yaozhong Shi, Zachary E Ross, Domniki Asimaki, and Kamyar Azizzadenesheli. Stochastic process learning via operator flow matching. In *Advances in Neural Information Processing Systems*, 2025.
- Dule Shu, Zijie Li, and Amir Barati Farimani. A physics-informed diffusion model for high-fidelity flow field reconstruction. *Journal of Computational Physics*, 478:111972, 2023.
- Matthew Tancik, Pratul P Srinivasan, Ben Mildenhall, Sara Fridovich-Keil, Nithin Raghavan, Utkarsh Singhal, Ravi Ramamoorthi, Jonathan T Barron, and Ren Ng. Fourier features let networks learn high frequency functions in low dimensional domains. In *Advances in Neural Information Processing Systems*, volume 33, pp. 7537–7547, 2020.
- Ricardo Vinuesa, Steven L Brunton, and Beverley J McKeon. The transformative potential of machine learning for experiments in fluid mechanics. *Nature Reviews Physics*, 5:536–545, 2023.
- Qing Wang, Scott Goodrick, Matthias Ihme, Yi-Fan Chung, John Anderson, Pierre Bénard, and et al. FireBench: A high-fidelity ensemble simulation framework for exploring wildfire behavior and data-driven modeling. *arXiv preprint arXiv:2406.08589*, 2024.
- Sifan Wang, Zhikai Wu, David van Dijk, and Lu Lu. Geometric function flow matching for inverse operator learning over complex geometries. *arXiv preprint arXiv:2509.24117*, 2025.
- Sifan Wang, Zehao Dou, Tong-Rui Shan, Yuji Liu, and Lu Lu. FunDiff: Diffusion models over function spaces for physics-informed generative modeling. *Nature Communications*, 17:5749, 2026.
- Haixu Wu, Huakun Luo, Haowen Wang, Jianmin Wang, and Mingsheng Long. Transolver: A fast transformer solver for PDEs on general geometries. In *International Conference on Machine Learning*, 2024.
- Wei jie Yang, Xun Zhang, Simin Jiang, and Yubao Zhou. PIS: A generalized physical inversion solver for arbitrary sparse observations via set conditioned flow matching. *arXiv preprint arXiv:2512.13732*, 2025.
- Jiachen Yao, Abbas Mammadov, Julius Berner, Gavin Kerrigan, Jong Chul Ye, Kamyar Azizzadenesheli, and Anima Anandkumar. Guided diffusion sampling on function spaces with applications to PDEs. In *Advances in Neural Information Processing Systems*, 2025.

Anthony Zhou, Zijie Li, Michael Schneier, John R Buchanan Jr, and Amir Barati Farimani. Text2PDE: Latent diffusion models for accessible physics simulation. In International Conference on Learning Representations, 2025.

Yilin Zhuang, Sibor Cheng, and Karthik Duraisamy. Spatially-aware diffusion models with cross-attention for global field reconstruction with sparse observations. Computer Methods in Applied Mechanics and Engineering, 435:117623, 2025.

A Architecture details

The velocity network used for all experiments has 6.51M trainable parameters (parameter-matched to the latent flow-matching baseline’s two stages combined). Sensor tokens and queries use Fourier positional encodings with 32 bands and maximum frequency 64. The global pathway compresses M sensor tokens into $S=128$ latents of width 256 with 8-head cross-attention, refined by 4 latent self-attention blocks with feed-forward multiplier 4; sensor re-injection cross-attention runs after every block, and a final back-attention pass writes latent context into the sensor tokens. The per-query readout attends from a coordinate-conditioned decoder token to the latent array. The local pathway gathers the $K=16$ nearest sensor tokens under the importance-warped metric of Section 3.2 with a learnable RBF bandwidth; a training-time ablation on a trained checkpoint found $K=64$, 32, and 16 indistinguishable in reconstruction error (relative L_2 0.273/0.273/0.272), so the smallest is used. Neighbor search and RBF aggregation run as streaming KeOps reductions (Charlier et al., 2021) with query chunking (2,048 queries per chunk for the global readout, 4,096 for decoding). The source GP uses 256 random Fourier features with correlation length 0.15 in normalized coordinates.

Two implementation details matter at our scale. First, cross-attention over tens of thousands of variably valid sensor tokens with a key-padding mask forces the memory-efficient (CUTLASS) attention backward, which at $M \approx 39k$ keys is $\sim 70\times$ slower at the kernel level than the flash backward; selecting each item’s valid keys by boolean mask and running unmasked attention per item is mathematically identical (max deviation 4×10^{-6}) and doubles training throughput ($0.88 \rightarrow 0.41$ s per step, compiled). Second, during ODE sampling, all sensor-branch activations, neighbor geometry, and the coordinate-conditioned readout are cached after the first step (Section 3.3); the cache is exact (bitwise-zero output difference) and reduces a 16-step 125^3 solve from 69.5 s to 5.1 s.

B Training details

All models train on single 80 GB H100 GPUs in bf16 autocast with an exponential moving average of weights (decay 0.9995) used for evaluation, Adam at learning rate 10^{-4} , weight decay 10^{-6} , and batch size 20 point-cloud samples. Turbulence: 6,000 epochs (~ 20 h wall-clock; the latent baseline’s two stages total ~ 21 h, so total budgets are matched), 39,062 query points per sample per step, per-channel sensor counts drawn log-uniformly between 0.1% and 1% of points, observed channels $\{U_x, U_z\}$, t logit-normal, $\lambda_{sp}=0.02$ with a 32^3 spectral block and 12 bins. Wildfire: 3,500 epochs, observed channels $\{u, v, w\}$, sensor counts 0.03–1%, operator amortization with noise amplitude up to 0.1σ , slab occlusion probability 0.2, and observed-channel dropout probability 0.3; lateral-reflection augmentation. Wing: 4,000 epochs, surface pool of 512 pressure + 128 shear sensors drawn per step from the 6.5×10^4 -node skin, Mach and incidence tokens, spanwise-reflection augmentation. The forward pass is compiled; compilation preserves checkpoint compatibility.

C Datasets and preprocessing

Isotropic turbulence. Four 125^3 cutouts at disjoint spatial locations of the JHTDB isotropic1024coarse DNS, 50 snapshots each at a stride of 100 database steps, fields (U_x, U_y, U_z, p) , per-field standardized. Training uses cutouts 1–3 (150 snapshots); all reported evaluation uses cutout 4. Within the training cutouts, temporal blocks are kept

contiguous. The leakage analysis motivating this protocol (Section 4) was run on 617 consecutive frames of a single cutout: frame-to-frame correlation is 1.00 at unit separation and 0.67 at separation 100, and a shuffled split placed every validation frame within one frame of a training frame.

SHIFT-WING. 673 converged RANS cases spanning the dataset’s Mach blocks (0.5–0.85), incidence 0.02–3.99°, and 40 geometry parameters (aspect ratio 7.5–11, sweep 25–37°). Each case is reduced to 4×10^5 near-body volume nodes (subsamped from the adaptive mesh) and the full 6.5×10^4 -node surface pool carrying C_p and the three wall-shear components as distinct observed-quantity indices; Mach and incidence enter as parameter tokens with dedicated indices. Split 600/73 by case; physical-coordinate spanwise reflection augments training.

FireBench. From the public FireBench store, the two LES runs at 10 and 12 m/s driving wind. The 3D extraction crops a $152 \times 126 \times 192$ region around the fire front (streamwise 250–706 m at 3 m, full lateral extent at 2 m, 0–192 m altitude at 1 m), avoiding the empty sky aloft, with fields $(u, v, w, \theta, \rho_f)$; 60 snapshots per run at stride 2, merged to 120, temporally blocked split with a 10-frame guard gap.

D Baseline configurations

Latent flow matching. A 3D convolutional autoencoder (3 levels, $\sim 16^3 \times 48$ latent) trained for 5,000 epochs on full fields, followed by 10,000 epochs of rectified flow in the latent space with a point-encoder conditioning branch over the identical sensor draws; combined parameter count 6.53M, matched to ours. Identical data, splits, augmentation, and sensor protocol. The two-stage structure gives this baseline roughly $50\times$ more supervised field values per step than LOCUS; the supervision ablation in Appendix F addresses the resulting confound.

Gen4Turb. The released code and architecture of Oommen et al. (2026) (elucidated-diffusion U-Net, channel dimension 16, self-conditioning, mask-channel conditioning), trained on our turbulence data at its native 120^3 crop of the 125^3 region, with σ_{data} taken from its own preprocessing. Its training script defaults to a compute budget an order of magnitude beyond every other method; we checkpoint every 10 epochs and report the wall-clock-matched checkpoint (epoch 4,930, matching our 20 h). Two observation protocols are evaluated: its native shared-mask protocol (all four channels observed at the same 1% locations; more information than any other method receives) and our strict protocol (two channels at independent locations). Because the strict protocol is out of distribution for the shared-mask training, we additionally retrained it under our exact observation model (independent per-channel masks, coverage drawn 0.1–1%); the retrain moved strict-protocol accuracy by less than 0.01 (0.700→0.706), and its validation-selected checkpoint is substantially worse than the budget-matched one (0.921 vs 0.706), so the table reports the budget-matched retrained model, the choice most generous to the baseline.

Senseiver. The attention-based deterministic reconstructor of Santos et al. (2023) under identical sensor randomization, augmentation, and splits, trained to its own convergence. As a deterministic model its CRPS equals its MAE, and no dispersion metrics apply.

Operator evaluations. All test-time operators (noise, slab occlusion, channel dropout) are applied to the assembled sensor sets by shared seeded code, so every method sees bit-identical corrupted observations per snapshot.

E Metric definitions

For an ensemble $\{s_k\}_{k=1}^K$ and truth y at a point, the fair CRPS estimator is $\text{CRPS} = \frac{1}{K} \sum_k |s_k - y| - \frac{1}{2K(K-1)} \sum_{k \neq l} |s_k - s_l|$, averaged over points and fields in standardized units. The spread–error ratio is the ratio of the ensemble standard deviation (RMS over

points) to the RMS error of the ensemble mean; 1.0 indicates dispersion consistent with error. Coverage₉₀ is the fraction of points whose truth falls inside the central 90% ensemble interval. Rank histograms accumulate the rank of the truth within the ordered ensemble per point. Relative L_2 is $\|\hat{u} - u\|_2 / \|u\|_2$ per field, reported for the ensemble mean and for individual samples.

F Additional results

Table 3: Spectral-weight ablation on the turbulence benchmark (matched protocol, $K=8$, 8 snapshots). Ratios are single-sample inertial-band spectral power relative to truth at NFE 2. The weight trades single-sample sharpness against posterior dispersion.

λ_{sp}	Rel. L_2 mean	Rel. L_2 sample (NFE 2)	CRPS	Spread/Err.	Cov. ₉₀	Inertial ratio
0	0.612	0.730	0.320	0.79	0.65	0.43
0.02	0.592	0.651	0.333	0.64	0.55	0.48
0.05	0.619	0.676	0.340	0.68	0.59	0.48

Table 4: Sampling-step (NFE) sweep for the $\lambda_{\text{sp}}=0$ model ($K=4$): dispersion improves monotonically with step count at flat CRPS; the conditioning cache makes the sweep nearly cost-free (4.4–5.1 s per 125^3 sample).

NFE	Rel. L_2 (mean)	CRPS	Spread/Error	Coverage ₉₀
2	0.607	0.299	0.65	0.41
4	0.635	0.294	0.82	0.51
8	0.653	0.294	0.89	0.54
16	0.662	0.295	0.92	0.55

Table 5: Wildfire per-field relative L_2 (ensemble mean) for LOCUS: cross-variable inference of unobserved thermochemical fields from wind observations.

Operator	u	v	w	θ (unobs.)	ρ_f (unobs.)
Clean (wind observed)	0.442	0.726	0.583	0.487	0.044
Channel dropout (u only)	0.470	1.026	0.932	0.524	0.064

Augmentation and supervision ablations (turbulence). With octahedral augmentation as the base (best held-out flow-matching loss 0.7475), adding translations and restriction to proper rotations (0.7515), continuous rotations (0.7783), or narrowing the latent array (0.7648) all fall within within-run fluctuation (0.05–0.11); no intervention improved on plain octahedral. Without any augmentation, all methods—including both baselines—diverge on the spatially disjoint region, confirming that at ~ 150 training snapshots symmetry augmentation is what makes cross-region generalization possible at all. Doubling the query budget to 78k points per step (three seeds: 0.756, 0.790, 0.767) does not improve on 39k (0.7475). A heavy-regularization arm (weight decay 10^{-2} , dropout 0.1) is likewise null (0.769–0.786), indicating the binding constraint is snapshot count, not overfitting control or per-step supervision.

Cost at the benchmark resolution. At 125^3 : LOCUS trains at 0.44s/step, 39.9GB, and samples a full field in 4.5s (NFE 4) at 1.4GB; the latent baseline trains its stages at 8.2/3.4GB and samples in 0.7s at 0.25GB; Senseiver trains at 0.17s/step, 19.4GB, and decodes in 0.13s at 1.2GB; Gen4Turb trains at 57.5GB (production configuration, 120^3) and samples in ~ 1 s (32 denoising steps). Below the $\sim 240^3$ crossover, grid computation is cheaper; the separation of Figure 3 is about what happens above it.

Table 6: Robust-vs-clean training ablation on the wildfire operator matrix (relative L_2 / CRPS, $K=8$, 8 snapshots). The clean-trained control saw no corrupted sensors during training.

Operator	LOCUS (operator-amortized)	LOCUS (clean-trained)
Clean	0.456 / 0.234	0.455 / 0.231
Noise $\sigma=0.1$	0.458 / 0.237	0.456 / 0.234
Noise $\sigma=0.3$	0.461 / 0.248	0.458 / 0.246
Slab occlusion (25%)	0.511 / 0.275	0.513 / 0.277
Channel dropout (u only)	0.603 / 0.355	0.743 / 0.436

Pending additions.